

# XGBoost Regression

XgBoost is an ensemble learner.

XGBoost regression or Extreme Gradient Boosting is supervised learning technique which uses Decision Trees and and Gradient Boosting method. It contains loss function and regularized parameters which gives ot name Extreme.

It tells about the difference between actual values and predicted values, i.e how far the model results are from the real values.

## Methodology

1. Import all the necessary libraries for the task.
2. Read Data into the DataFrame.
3. Understand Data by using methods like hea() and info().
4. Cleaning Data or Preprocessing on DataSet.
5. Visualizing Data using Scatterplots, histograms and Heatmaps.
6. Splitting Data into features and Target varieable.
7. Splitting Data into Training and Testing Data.
8. Training model using XGBoost Regressor on training Data.
9. Predicting values on Testing Data.
10. Calculatin RMSE and R2 Score on Testing Data to see how well model preformed.m

## Importing Libraries

```
In [1]:

import pandas as pd
from sklearn.metrics import mean_squared_error, r2_score
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
import xgboost as xgb
%matplotlib inline
```

## Reading Dataset From .csv file as Pd Data Frame

```
In [2]:

DataSet = pd.read_csv("Asteroid_Updated.csv")

C:\Users\nabhr\Anaconda3\lib\site-packages\IPython\core\interactiveshell.py:3057: DtypeWarning: Columns (0,10,15,16,23,24) have mixed types. Specify dtype option on import or set low_memory=False.
  interactivity=interactivity, compiler=compiler, result=result)
```

## Printing first 10 Rows of Data Set to cunderstand the data.

```
In [3]:

DataSet.head(10)
```

Out[3]:

|   | name    | a        | e        | i         | om         | w          | q        | ad       | per_y    | data_arc | ... | UB    | IR  | spec_B | spec_T | G    | moid     | class | n        | p          |
|---|---------|----------|----------|-----------|------------|------------|----------|----------|----------|----------|-----|-------|-----|--------|--------|------|----------|-------|----------|------------|
| 0 | Ceres   | 2.769165 | 0.076009 | 10.594067 | 80.305532  | 73.597694  | 2.558684 | 2.979647 | 4.608202 | 8822.0   | ... | 0.426 | NaN | C      | G      | 0.12 | 1.594780 | MBA   | 0.213885 | 1683.14571 |
| 1 | Pallas  | 2.772466 | 0.230337 | 34.836234 | 173.080063 | 310.048857 | 2.133865 | 3.411067 | 4.616444 | 72318.0  | ... | 0.284 | NaN | B      | B      | 0.11 | 1.233240 | MBA   | 0.213503 | 1686.15591 |
| 2 | Juno    | 2.669150 | 0.256942 | 12.988919 | 169.852760 | 248.138626 | 1.983332 | 3.354967 | 4.360814 | 72684.0  | ... | 0.433 | NaN | Sk     | S      | 0.32 | 1.034540 | MBA   | 0.226019 | 1592.78721 |
| 3 | Vesta   | 2.361418 | 0.088721 | 7.141771  | 103.810804 | 150.728541 | 2.151909 | 2.570926 | 3.628837 | 24288.0  | ... | 0.492 | NaN | V      | V      | 0.32 | 1.139480 | MBA   | 0.271609 | 1325.43271 |
| 4 | Astraea | 2.574249 | 0.191095 | 5.366988  | 141.576605 | 358.687607 | 2.082324 | 3.066174 | 4.130323 | 63507.0  | ... | 0.411 | NaN | S      | S      | NaN  | 1.095890 | MBA   | 0.238632 | 1508.60041 |
| 5 | Hebe    | 2.425160 | 0.203007 | 14.737901 | 138.640203 | 239.807490 | 1.932835 | 2.917485 | 3.776755 | 62329.0  | ... | 0.399 | NaN | S      | S      | 0.24 | 0.973965 | MBA   | 0.260972 | 1379.45971 |
| 6 | Iris    | 2.385334 | 0.231206 | 5.523651  | 259.563231 | 145.265106 | 1.833831 | 2.936837 | 3.684105 | 62452.0  | ... | 0.484 | NaN | S      | S      | NaN  | 0.846100 | MBA   | 0.267535 | 1345.61911 |
| 7 | Flora   | 2.201764 | 0.156499 | 5.886955  | 110.889330 | 285.287462 | 1.857190 | 2.546339 | 3.267115 | 62655.0  | ... | 0.489 | NaN | NaN    | S      | 0.28 | 0.874176 | MBA   | 0.301681 | 1193.31371 |
| 8 | Metis   | 2.385637 | 0.123114 | 5.576816  | 68.908577  | 6.417369   | 2.091931 | 2.679342 | 3.684806 | 61821.0  | ... | 0.496 | NaN | NaN    | S      | 0.17 | 1.106910 | MBA   | 0.267484 | 1345.87531 |
| 9 | Hygiea  | 3.141539 | 0.112461 | 3.831560  | 283.202167 | 312.315206 | 2.788240 | 3.494839 | 5.568291 | 62175.0  | ... | 0.351 | NaN | C      | C      | NaN  | 1.778390 | MBA   | 0.177007 | 2033.81821 |

10 rows x 31 columns

# Analysing Data Coloumns and it's types

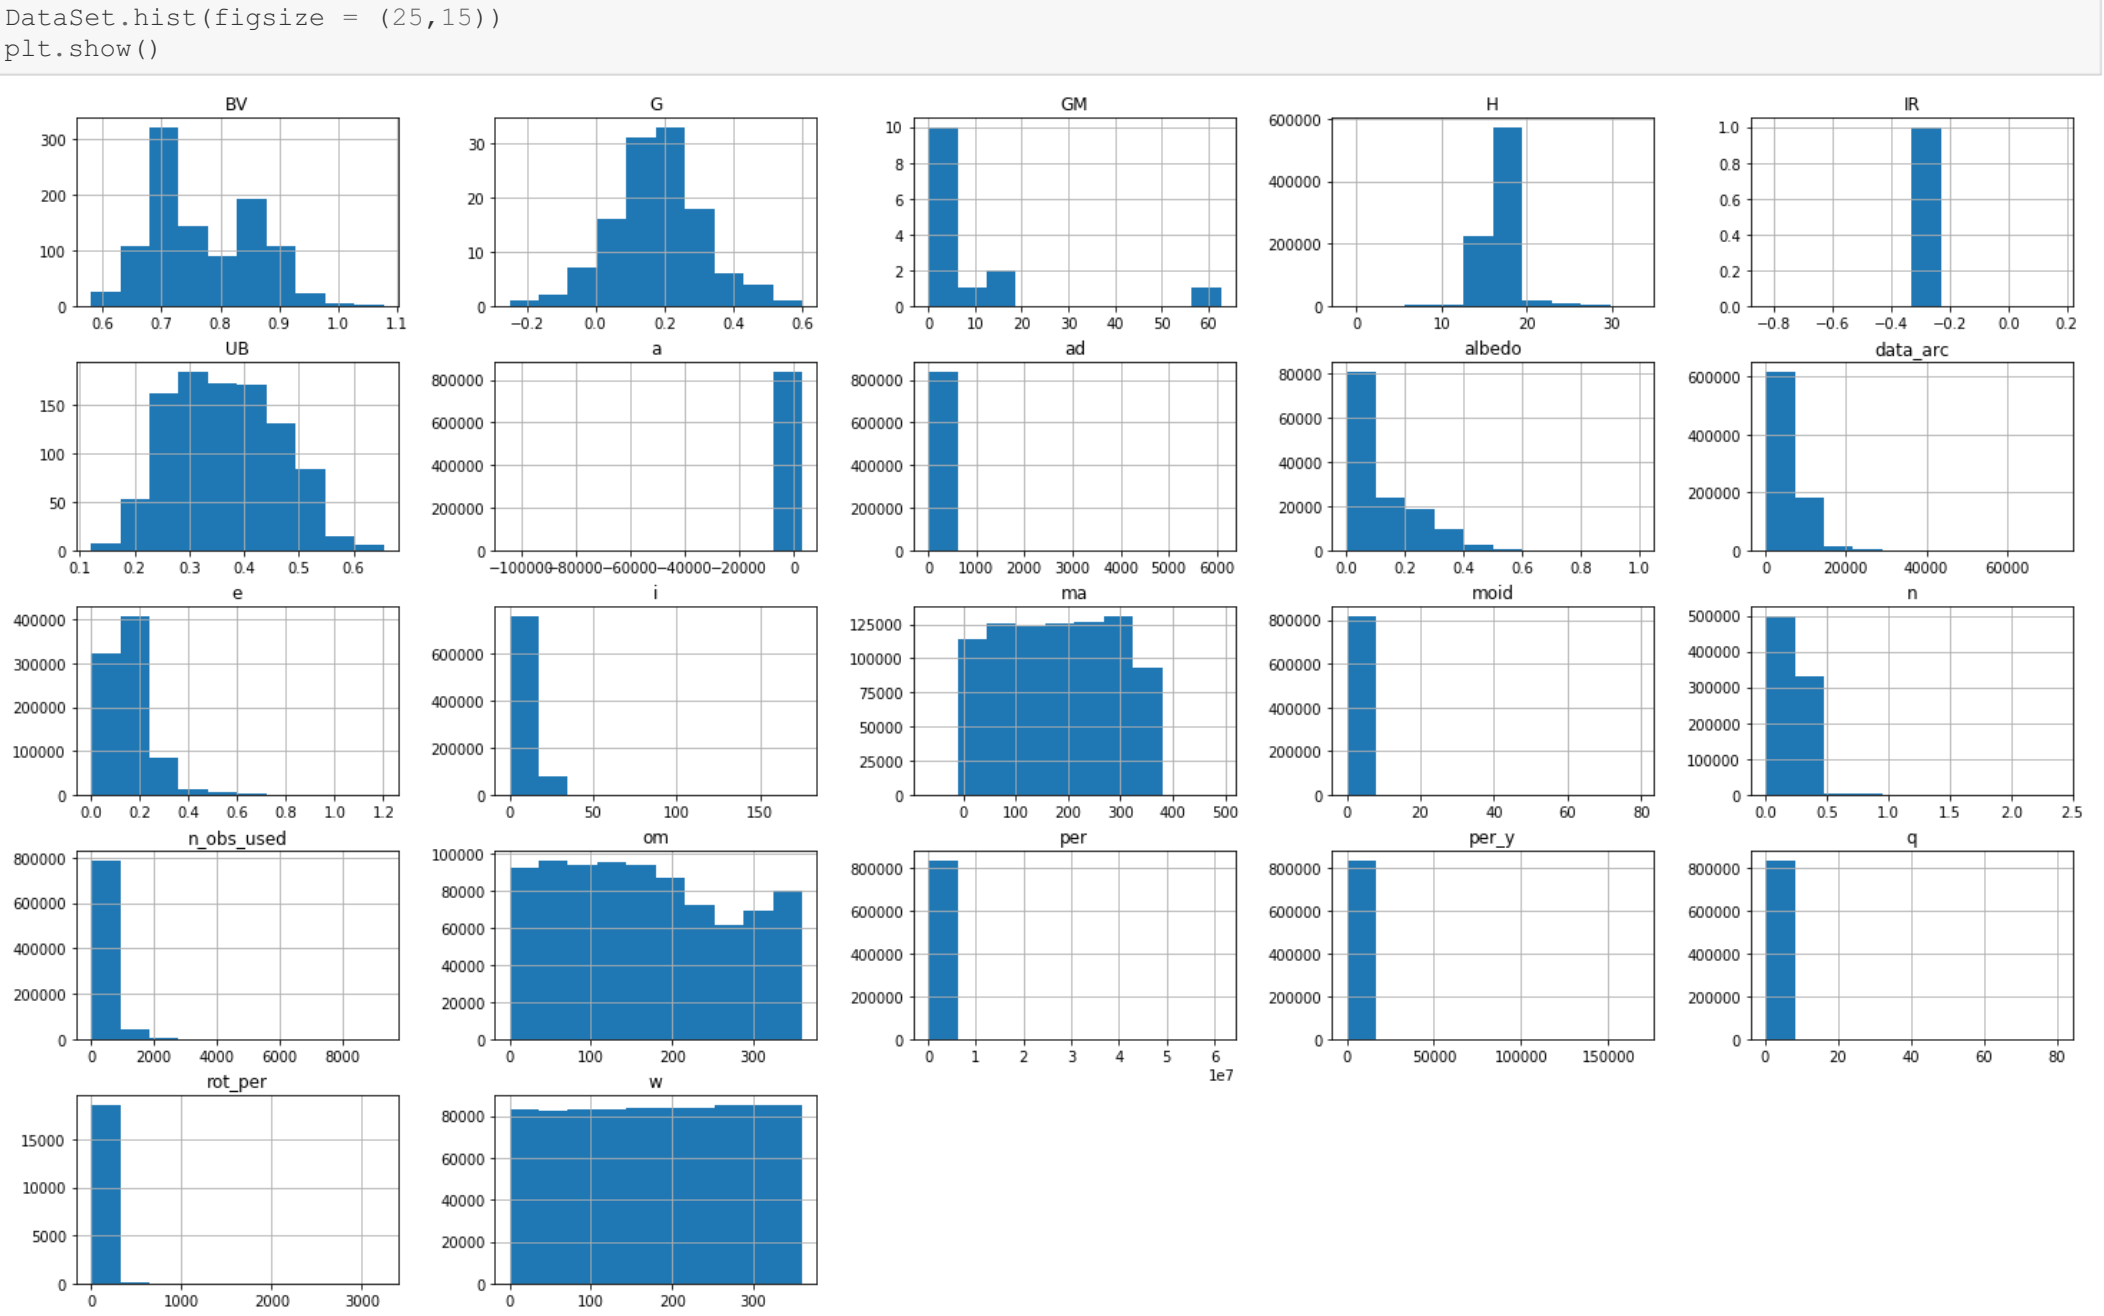
In [4]:

```
DataSet.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 839714 entries, 0 to 839713
Data columns (total 31 columns):
name                21967 non-null object
a                   839712 non-null float64
e                   839714 non-null float64
i                   839714 non-null float64
om                  839714 non-null float64
w                   839714 non-null float64
q                   839714 non-null float64
ad                  839708 non-null float64
per_y               839713 non-null float64
data_arc            824240 non-null float64
condition_code      838847 non-null object
n_obs_used          839714 non-null int64
H                   837025 non-null float64
neo                 839708 non-null object
pha                 823272 non-null object
diameter            137636 non-null object
extent              18 non-null object
albedo              136409 non-null float64
rot_per             18796 non-null float64
GM                  14 non-null float64
BV                  1021 non-null float64
UB                  979 non-null float64
IR                  1 non-null float64
spec_B              1666 non-null object
spec_T              980 non-null object
G                   119 non-null float64
moid                823272 non-null float64
class               839714 non-null object
n                   839712 non-null float64
per                 839708 non-null float64
ma                  839706 non-null float64
dtypes: float64(21), int64(1), object(9)
memory usage: 198.6+ MB
```

## Visualizing Data and analyzing for irregularities

In [5]:



Since there are many coloumns which have very less non-null values. We can drop all those coloumns which have non-null values much less than total number of values.

In [6]:

```
DataSet = DataSet.drop('IR', axis = 1)
```

```
DataSet = DataSet.drop('UB', axis = 1)
DataSet = DataSet.drop('G', axis = 1)
DataSet = DataSet.drop('GM', axis = 1)
DataSet = DataSet.drop('extent', axis = 1)
DataSet = DataSet.drop('BV', axis = 1)
DataSet = DataSet.drop('spec_B', axis = 1)
DataSet = DataSet.drop('spec_T', axis = 1)
DataSet = DataSet.drop('rot_per', axis = 1)
DataSet = DataSet.drop('name', axis = 1)
```

**Analysing Data Coloumn Again and diving deep itno data cleaning and processing**

In [7]:

```
DataSet.shape
DataSet.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 839714 entries, 0 to 839713
Data columns (total 21 columns):
a                839712 non-null float64
e                839714 non-null float64
i                839714 non-null float64
om              839714 non-null float64
w               839714 non-null float64
q               839714 non-null float64
ad              839708 non-null float64
per_y           839713 non-null float64
data_arc        824240 non-null float64
condition_code  838847 non-null object
n_obs_used      839714 non-null int64
H               837025 non-null float64
neo             839708 non-null object
pha             823272 non-null object
diameter        137636 non-null object
albedo          136409 non-null float64
moid            823272 non-null float64
class           839714 non-null object
n               839712 non-null float64
per             839708 non-null float64
ma              839706 non-null float64
dtypes: float64(15), int64(1), object(5)
memory usage: 134.5+ MB
```

**As analysed from info, Dataset still has rows with null entry, se lets Delete those rows which have any entries as NA values.**

In [8]:

```
DataSet = DataSet.dropna(axis = 'index', how = 'any')
```

In [9]:

```
DataSet.shape
```

Out[9]:

```
(136005, 21)
```

**Now checking the rows with object values.**

In [10]:

```
for col in DataSet.columns:
    if(DataSet[col].dtype == object):
        A = DataSet[col].unique()
        print(col, A)
```

```
condition_code [0 1 3 2 '0' '1' '2' '4' '5' '9' '3' '7' 5.0 6.0 4.0 7.0 9.0 8.0 '8' '6']
neo ['N' 'Y']
pha ['N' 'Y']
diameter ['939.4' '545' '246.596' ... 0.122 0.6509999999999999 1.077]
class ['MBA' 'OMB' 'MCA' 'AMO' 'IMB' 'TJN' 'CEN' 'APO' 'ATE' 'AST' 'TNO']
```

**Dropping those columns with very large number of unique values or very small number of unique values**

In [11]:

```
DataSet = DataSet.drop('condition_code', axis = 1)
DataSet = DataSet.drop('neo', axis = 1)
DataSet = DataSet.drop('pha', axis = 1)
DataSet = DataSet.drop('class', axis = 1)
```

**Since diameter column only has in values as strings, lets convert diameter and n\_obs\_used column to float.**

In [12]:

```
DataSet['diameter'] = DataSet['diameter'].astype(float)
DataSet['n_obs_used'] = DataSet['n_obs_used'].astype(float)
```

In [13]:

```
DataSet.info()
```

DataSet.shape

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 136005 entries, 0 to 810375
Data columns (total 17 columns):
a                136005 non-null float64
e                136005 non-null float64
i                136005 non-null float64
om              136005 non-null float64
w                136005 non-null float64
q                136005 non-null float64
ad              136005 non-null float64
per_y           136005 non-null float64
data_arc        136005 non-null float64
n_obs_used      136005 non-null float64
H               136005 non-null float64
diameter         136005 non-null float64
albedo          136005 non-null float64
moid            136005 non-null float64
n               136005 non-null float64
per             136005 non-null float64
ma              136005 non-null float64
dtypes: float64(17)
memory usage: 18.7 MB
```





Out[13]:







(136005, 17)


### Plotting Scatter Plots.

In [14]:

```
for col in DataSet.columns:
    plt.scatter(DataSet.loc[:,col], DataSet.loc[:, 'diameter'], color = 'black')
    plt.legend([col],loc = 'upper right')
    plt.show()
```








Plotting Heatmap to better visualize the features."

In [15]:

```
def Heat(Datax):
    correlate = Datax.corr()
    fig, ax = plt.subplots(figsize = (20,10))
    ax = sns.heatmap(correlate, annot = True)
    plt.show()
Heat(DataSet)
```



Since XGBoost used Decision Trees which checks for every features individually in stumps to train the model, we don't need to eliminate any features based on coreltion matrix.

In [16]:

```
DataSet.info()
DataSet.shape
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 136005 entries, 0 to 810375
Data columns (total 17 columns):
a                136005 non-null float64
e                136005 non-null float64
i                136005 non-null float64
om              136005 non-null float64
w               136005 non-null float64
q               136005 non-null float64
ad              136005 non-null float64
per_y           136005 non-null float64
data_arc        136005 non-null float64
n_obs_used      136005 non-null float64
H               136005 non-null float64
diameter        136005 non-null float64
albedo          136005 non-null float64
moid            136005 non-null float64
n               136005 non-null float64
per             136005 non-null float64
ma              136005 non-null float64
dtypes: float64(17)
memory usage: 18.7 MB
```

Out[16]:

(136005, 17)

### Splitting Data Set into Features and Target Variables.

In [17]:

```
Y = DataSet.loc[:, 'diameter']
X = DataSet.drop('diameter', axis = 1)
```

### Splitting Data Sets into Training and Testing Data

In [18]:

```
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.25, random_state=5)
```

### Applying XGBoost to train a model 'XG'.

In [19]:

```
XG = xgb.XGBRegressor(objective='reg:linear', colsample_bytree = 0.5, learning_rate = 0.1,
                        max_depth = 5, alpha = 10, n_estimators = 20)
```

Fitting our training data to the model

```
In [20]:
XG.fit(X_train,y_train)

[12:51:58] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.2.0/src/objective/regression_obj.cu:174: reg:linear is now deprecated in favor of reg:squarederror.
[12:51:59] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.2.0/src/objective/regression_obj.cu:174: reg:linear is now deprecated in favor of reg:squarederror.

Out[20]:
XGBRegressor(alpha=10, base_score=0.5, booster='gbtree', colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=0.5, gamma=0, gpu_id=-1,
              importance_type='gain', interaction_constraints='',
              learning_rate=0.1, max_delta_step=0, max_depth=5,
              min_child_weight=1, missing=nan, monotone_constraints='()',
              n_estimators=20, n_jobs=0, num_parallel_tree=1,
              objective='reg:linear', random_state=0, reg_alpha=10, reg_lambda=1,
              scale_pos_weight=1, subsample=1, tree_method='exact',
              validate_parameters=1, verbosity=None)
```

Predicting values of testing Data using the model.

```
In [21]:
Y_Predict = XG.predict(X_test)
```

Calculating RMSE and r2 Score based on predictions done on Testing Data

```
In [22]:
rmse = np.sqrt(mean_squared_error(y_test, Y_Predict))
print("RMSE: %f" % (rmse))
r2 = r2_score(y_test, Y_Predict)
print(r2)

RMSE: 2.213173
0.9274388301899217
```

RMSE for predicted values on Testing Data: 2.21.

R2 Score : .92

```
In [ ]:
```